

Matrix Exponential Optimization

Advanced Systems Lab

Team 28



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Algorithm

■ Goal

- Algorithm 5.1 proposes a new way to approximate e^A , where A is a square matrix.

■ Motivation

- Previous implementation had poor precision with big matrices (overscaling)

■ Extensive use of matrix multiplication

- Wipes out cache
- *Blas* allowed for MMM
- Main focus in optimization code fragments between MMM

Infrastructure

■ Validation

- Existing implementation: *scipy expm* in `scipy.linalg` module (Python)
- The C++ infrastructure interacts with a Python script to get the a valid computation

■ Timing

- Processor used: i5-6400, Skylake microarchitecture @2700 Mhz
- Cycles computed using *TSC* counter

Cost analysis

■ Cost measure

- Total number of FLOPS
- All floating-point operations have the same weight
- Code instructed to count FLOPS through macros

```
#ifdef FLOPS
#define FLOPS_RESET
#define FLOPS_ADD(x)
#define FLOPS_ADD_N(x)
#define FLOPS_ADD_N2(x)
#define FLOPS_ADD_N3(x)
#define FLOPS_PRINT
#endif
```

Straightforward C Implementation

■ Codebase

- **mexp_basic.c**: performs the Matrix Exponential (Memexp), utilizing functions from **utils.c** and queries *normest* from **norms.c** to get an estimation of the norm from a matrix
- **utils.c** and **utils.h**: utility functions needed in **mexp_basic.c**, promoting modularity
- **norms.c**: houses *normest* for norm estimation, used by **mexp_basic.c** for matrix calculations

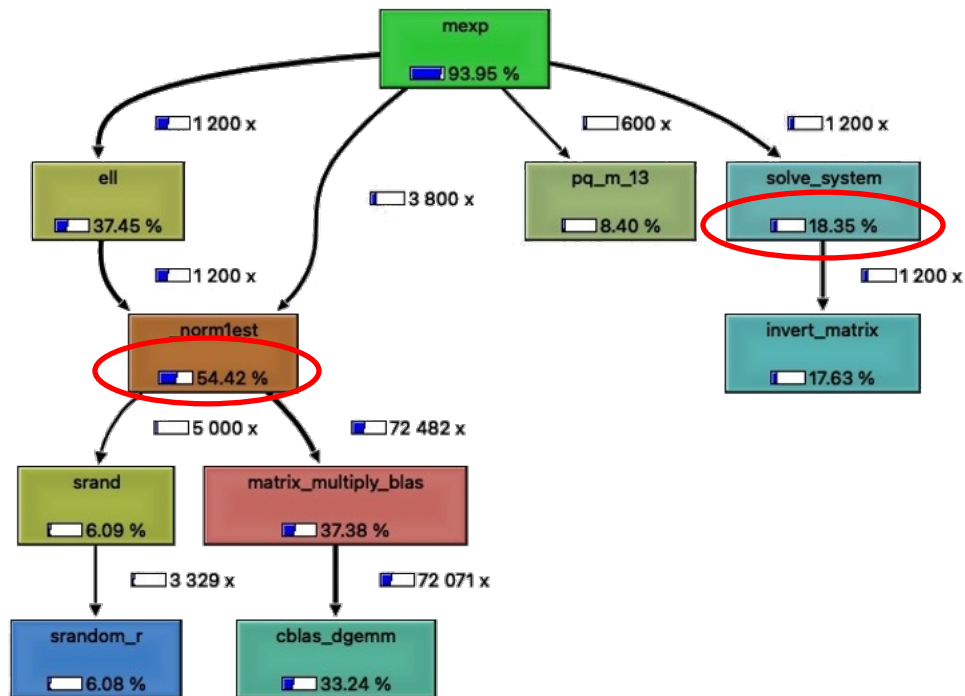
■ BLAS is used to perform MMMs

■ Matrices are represented as array in row-major order

Callgrind analysis

■ Bottlenecks

- solve-system
- normest




Basic optimizations

- **OPT-1 and OPT-2 modifications involved the following scalar optimizations:**
 - Ensure the constant use of arrays
 - Save precomputed data and reused after
 - Removed calls to `abs()`, `ceil()` or divisions
 - ILP improvements, unrolling and scalar replacement
 - Function inlining

Basic optimizations


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

BLAS



Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


no BLAS



Blocking Plot


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

BLAS



Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


no BLAS



Vectorization Plot


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

BLAS

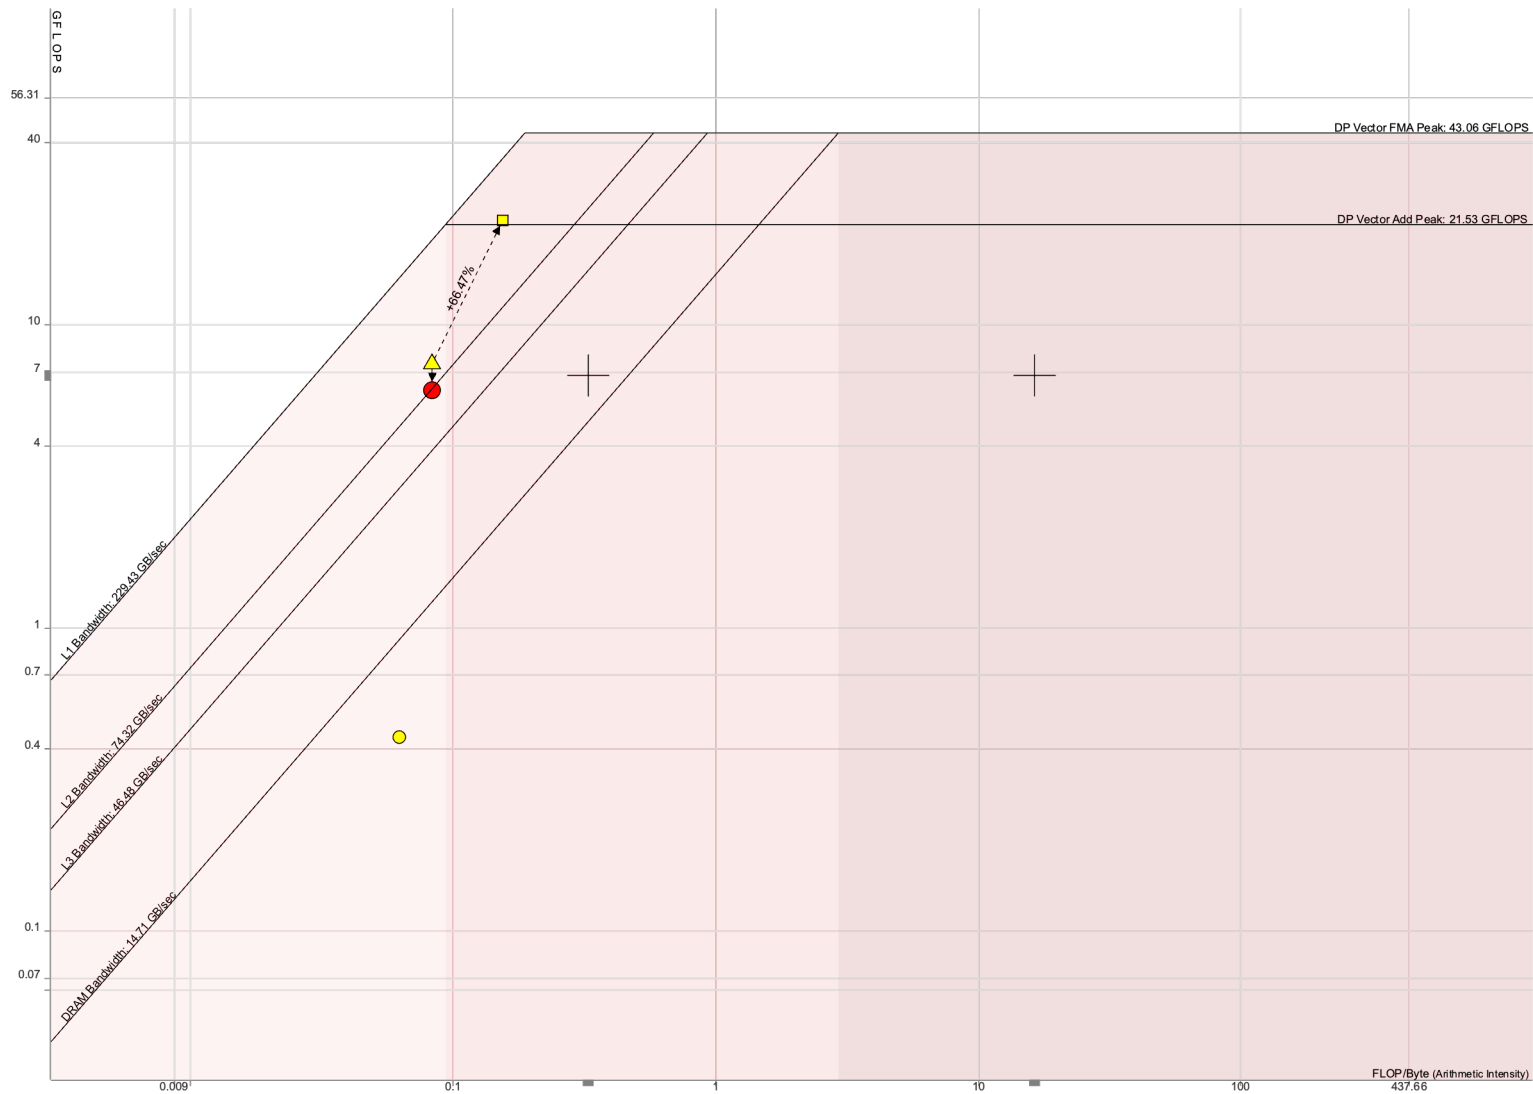


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

no BLAS



Roofline plot *inverse* (Base, BB and BV)



Gaussian Elimination

- The straightforward implementation used Gaussian Elimination to solve the system $QX=P$, where all matrices are $N \times N$.
- **Limitation**
 - No possible blocking in Gaussian Elimination
 - LU decomposition is the matrix form of Gaussian Elimination
- In LU decomposition computations can be reused by backward and forward substitution applied to every column of P .

Comparison

■ Gaussian Elimination optimizations

- Basic optimizations
 - *Inline swap_rows and avoid repeated calculations*
 - *Strength reduction*
- Vectorization

■ LU decomposition optimizations


- Basic optimizations
 - *Loop reordering provided a major improvement (x10)*
- Blocking
- Vectorization

Loop ordering

■ Improved spatial and temporal locality

```
// Forward substitution
for (int i = 0; i < n; i++){
    for (int j = 0; j < n; j++){
        for (int k = 0; k < j; k++){
            P[j*n + i] -= P[k*n + i] * Q[j*n + k];
        }
    }
}

// Backward substitution
for (int i = 0; i < n; i++){
    for (int j = n - 1; j >= 0; j--){
        double temp = P[j*n + i];
        for (int k = j + 1; k < n; k++){
            temp -= P[k*n + i] * Q[j*n + k];
        }
        P[j*n + i] = temp / Q[j*n + j];
    }
}
```



```
// Forward substitution
for (int j = 0; j < n; j++){
    for (int k = 0; k < j; k++){
        for (int i = 0; i < n; i++){
            P[j*n + i] -= P[k*n + i] * Q[j*n + k];
        }
    }
}

// Backward substitution
for (int i = n - 1; i >= 0; i--) {
    for (int k = n - 1; k > i; k--) {
        for (int j = 0; j < n; j++) {
            P[i*n + j] -= Q[i*n + k] * P[k*n + j];
        }
    }
    for (int j = 0; j < n; j++) {
        P[i*n + j] /= Q[i*n + i];
    }
}
```


Gaussian Elimination Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz

L1:64KB L2:256KB L3:6MB

Compiler: gcc 10.2.1

Performance [Flops/Cycle]




Gaussian Elimination Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz

L1:64KB L2:256KB L3:6MB

Compiler: gcc 10.2.1

Performance [Flops/Cycle]




Gaussian Elimination Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz

L1:64KB L2:256KB L3:6MB

Compiler: gcc 10.2.1

Performance [Flops/Cycle]




LU Decomposition Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz

L1:64KB L2:256KB L3:6MB

Compiler: gcc 10.2.1


Performance [Flops/Cycle]



LU Decomposition Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


Performance [Flops/Cycle]



LU Decomposition Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


Performance [Flops/Cycle]



LU Decomposition Performance Plot

Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


Performance [Flops/Cycle]



Runtime LU vs GE

Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


Runtime [Millisecond]



Mexp-LU - Base


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

BLAS



Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1


no BLAS



Mexp-LU - Vectorised


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

BLAS

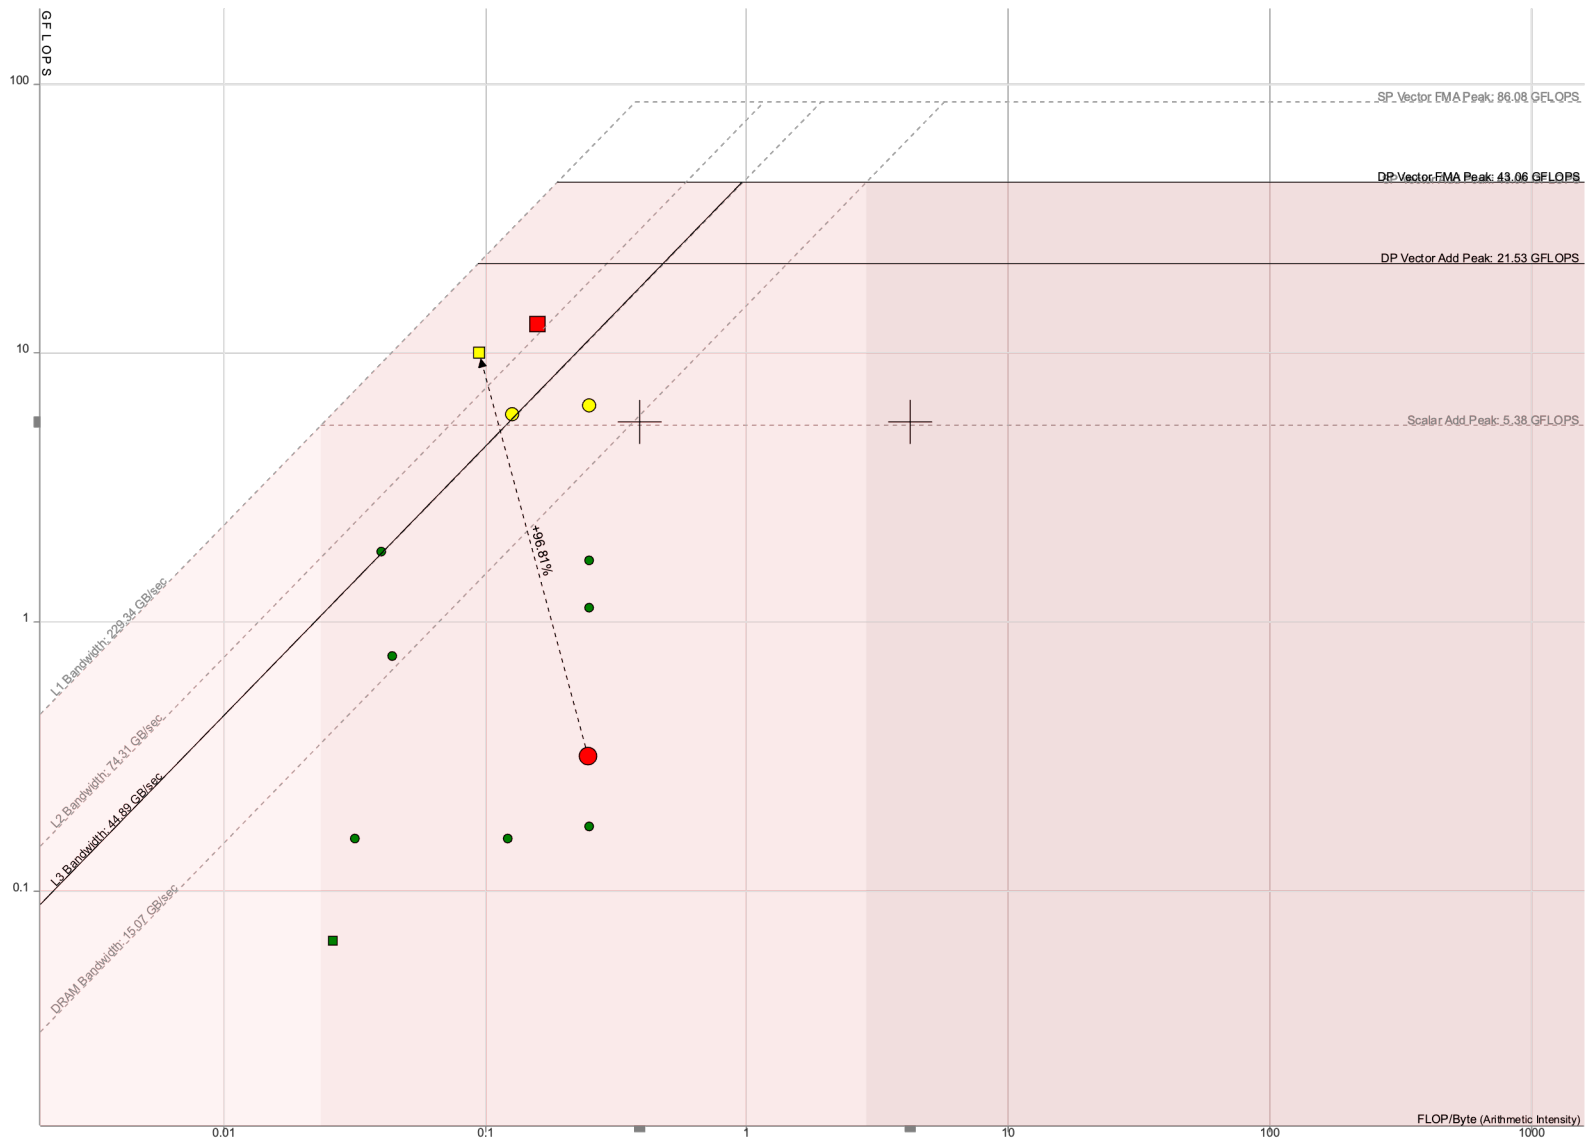


Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1

no BLAS




Roofline plot LU




Benchmark runtime comparison

Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1



Callgrind analysis

- Bottlenecks
 - solve-system
 - normest



Loop ordering

- Improved spatial and temporal locality

```
// Forward substitution
for (int i = 0; i < n; i++){
    for (int j = 0; j < n; j++){
        for (int k = 0; k < n; k++){
            P[i*n + j] = P[i*n + j] + Q[j*n + k];
        }
    }
}

// Backward substitution
for (int i = 0; i < n; i++){
    for (int j = n - 1; j >= 0; j--){
        double temp = P[i*n + j];
        for (int k = j + 1; k < n; k++){
            temp -= P[i*n + k] * Q[j*n + k];
        }
        P[i*n + j] = temp / Q[j*n + j];
    }
}
```




```
// Forward substitution
for (int j = 0; j < n; j++){
    for (int k = 0; k < n; k++){
        for (int i = 0; i < n; i++){
            P[i*n + j] = P[i*n + j] + Q[j*n + k];
        }
    }
}

// Backward substitution
for (int i = n - 1; i >= 0; i--){
    for (int k = n - 1; k >= 0; k--){
        for (int j = 0; j < n; j++){
            P[i*n + j] = P[i*n + j] + Q[j*n + k] * P[i*n + j];
        }
    }
    for (int j = 0; j < n; j++){
        P[i*n + j] = P[i*n + j] / Q[j*n + j];
    }
}
```

Q & A

Runtime LU vs GE

Intel(R) Core(TM) i5-6400 @ 2700Mhz
L1:64KB L2:256KB L3:6MB
Compiler: gcc 10.2.1



Roofline plot LU

